

Shivaum Shekhar Sharma

+91-7297088881 | sshivaum@gmail.com | github.com/shivaumsharma | linkedin.com/in/shivaum-sharma | Bengaluru

EDUCATION

Manipal Institute of Technology, Bengaluru | *B.Tech, Computer Science & Engineering (Data Science)* | 2023 - 2027

Relevant Coursework: Machine Learning, Natural Language Processing, Big Data Analytics, Database Management

PROJECTS

DigiNyaya - AI-Native Multi-Agent Dispute Resolution Platform | *diginyaya-frontend.onrender.com* |

github.com/shivaumsharma/DigiNyaya

- Built a RAG pipeline over 217 legal precedents with citation verification, so hallucinated citations are blocked before reaching users; 54/54 citations grounded across a repeatable 18-case golden-eval suite.
- Built grounding checks that drop any AI-generated document discrepancy unless its cited snippet is found verbatim in the source text; every flag carries its source snippet and confidence score for reviewers, and low-confidence cases escalate to human review (34.7% of 23,841 real court cases escalated).
- Evaluated on 23,841 real Indian court cases with an LLM judge: 74.9% directional accuracy and 39.4% full match on resolved cases, macro F1 0.63 vs. 0.45 for an always-claimant baseline. Full match fell from 50.8% on the first 1,914 cases, so earlier figures were retired. Found a judge-prompt bug that understated accuracy; reverted two fixes for a pro-landlord bias in tenancy cases that made real cases worse.
- Gated every PR with GitHub Actions CI (lint, tests, golden-eval suite); tracked token cost. Sarvam AI Startup Program (2026).

FinSight AI - Agentic Equity Research Platform | *web-ten-blond-39.vercel.app* | *github.com/shivaumsharma/FinSight-AI*

- Built a point-in-time backtesting framework over ~1,000 tickers (S&P 500/400/600); validation showed a prior "70% accuracy" figure only measured agreement with analyst consensus, so it was replaced with 12-month forward directional accuracy vs. an Always-Buy baseline.
- Caught 2 look-ahead data leaks and a class-imbalance trap in an ML classifier (61% raw accuracy was majority-class guessing).
- Added Snowflake as an optional analytics store for the backtest results: an explicit-schema table loaded via `write_pandas`, three SQL validation queries (precision by sector, Buy vs. Sell, accuracy trend by run date), and a degrade-to-no-op connection. It reproduced the published accuracy figures across 1,907 scored calls, with SQLite staying the system of record.
- Scored every RAG report on grounding and citation before display; handled edge cases (negative cash flow, missing line items) with a no-valuation fallback.
- Built with Next.js, FastAPI, LLM tool-calling, and a DCF engine encoding valuation business logic; used AI-assisted development (Claude Code, Codex, GitHub Copilot), verifying output with 1,000+ pytest tests in CI and a 65-finding code audit (including an RCE fix).

Amazon Financial Intelligence Platform | *Python, pandas, Streamlit, SQL (MySQL), scikit-learn* | *Streamlit Community Cloud*

- Built a Streamlit/Python and SQL analytics platform on Amazon's 2020-2024 10-K financials; found the FY2024 row held placeholder values and reconciled it against the actual earnings release and cash-flow statement.
- Implemented z-score anomaly detection in both pandas and a SQL window-function query, flagging FY2024 operating margin as anomalous.

EXPERIENCE

Intern, Advisory - Intelligent Business Systems | *PricewaterhouseCoopers (PwC)* | *India* | *Jun 2026 - Aug 2026*

- Authored 22 SOPs defining functionality, governance rules, and safety guardrails for a client's fintech AI systems; researched India's UPI transaction flow and RBI/NPCI requirements to keep new features compliant.
- Assessed agentic-AI workflows and use cases for a Fortune 500 e-commerce client and prototyped a small RAG agent; client approved all 10+ use cases, prioritizing 3 for Q3 2026.

Data Analyst Intern | *Department of Telecommunications, Govt. of India* | *Bhopal, India* | *May 2024 - Jun 2024*

- Analyzed a national telecom program's rollout status, budget allocation, and capacity gaps with SQL, Python, and Tableau; presented findings to senior IAS officers, who used them to guide funding decisions (covered in a Times of India article).

IT Intern | *Bharat Heavy Electricals Limited (BHEL)* | *Bhopal, India* | *May 2025 - Jun 2025*

- Built a Python graph engine for infrastructure dependency modeling with 100+ unit tests; cut a bottleneck from 40s to 0.1s (400x) on 10K-50K node graphs.
- Designed a REST API (FastAPI + PostgreSQL) for cascade-failure simulation and criticality scoring; scores reached 94% alignment with domain-expert rankings.

LEADERSHIP

President, MAHE Innovation Centre | *Bengaluru* | *Aug 2025 - Present*

- Lead 8-9 startups across 8 cells; mentored 12+ founders (INR 4L+ raised, 2 startups selected for an IIT Delhi cohort).

TECHNICAL SKILLS

Data & Validation: Snowflake, SQL (PostgreSQL, MySQL), Python, pandas, Streamlit, data reconciliation, data-quality checks, anomaly detection, Tableau, Power BI

AI Evaluation: explainable AI outputs (source citations, confidence scores, audit trail), golden-eval CI gates, citation-grounding checks, LLM-as-judge, retrieval evals (NDCG, MRR), confidence-driven escalation, backtesting, leakage checks

AI & Engineering: RAG, prompt engineering, LLM tool-calling, FastAPI, Docker, GitHub Actions, AWS, Terraform, Git, TypeScript